



PROGRAMMING ASSIGNMENT-02

Intelligent system

Recipe recommendation FastAPI Application

ITBIN-2110-0094 - Lakshitha Naveen

Contents

Introduction	2
Overview	2
Model Training.....	2
Data Collection.....	2
Model Architecture.....	2
Training Process	2
Evaluation	4
FastAPI Application	5
Overview	5
Application Structure	5
API Endpoints.....	5
Running the Application Locally.....	9
Deployment Setup	10
Railway Deployment - FastAPI - Swagger UI (railway.app).....	10
Deployment Steps.....	10
CI/CD Pipeline Documentation	11
CI/CD Pipeline Setup	11
Explanation of Steps	12
Online Testing	13
Conclusion.....	14
Summary	14

Introduction

Overview

This project provides a recipe recommendation system that takes a list of ingredients as input and returns a list of recommended recipes. The recommendation system uses a TF-IDF vectorizer for feature extraction and cosine similarity for matching.

Application link- [FastAPI - Swagger UI \(railway.app\)](#)

Github link - <https://github.com/LakshithaNaveenRathnasiri/FastAPI-application.git>

Model Training

Data Collection

Dataset: The dataset is a JSON file containing recipe data with ingredients and cuisine information. It is loaded from /content/recipe/train.json.

- Used Kaggle dataset.

Model Architecture

TF-IDF Vectorizer: Used for transforming ingredient texts into numerical feature vectors.

Cosine Similarity: Measures the similarity between the user's input ingredients and the recipes' ingredients.

Training Process

Data Preprocessing-

Ingredients are cleaned and tokenized.

Stopwords are removed and words are lemmatized.

Vectorization: The cleaned ingredients are vectorized using a TF-IDF Vectorizer with n-grams.

Recommendation Function: Calculates similarity scores and returns top recommendations based on user input.

```

code - import pandas as pd

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.model_selection import train_test_split

# Load the dataset
train_data_path = '/content/recipe/train.json'
train_df = pd.read_json(train_data_path)

# Clean and preprocess data
def preprocess_ingredients(ingredients):
    lemmatizer = WordNetLemmatizer()
    custom_stopwords = set(stopwords.words('english')).union({'fresh',
'chopped', 'sliced', 'diced', 'large', 'small', 'medium', 'extra', 'extra-
virgin', 'virgin'})

    cleaned_ingredients = []
    for ingredient in ingredients:
        ingredient = ingredient.lower()
        ingredient = re.sub(r'^a-z\s', '', ingredient)
        tokens = [lemmatizer.lemmatize(word) for word in ingredient.split() if
word not in custom_stopwords]
        cleaned_ingredients.append(' '.join(tokens))

    return ' '.join(cleaned_ingredients)

train_df['cleaned_ingredients'] =
train_df['ingredients'].apply(preprocess_ingredients)

# Vectorization with n-grams
vectorizer = TfidfVectorizer(stop_words='english', ngram_range=(1, 2))
ingredient_matrix = vectorizer.fit_transform(train_df['cleaned_ingredients'])

# Recommendation function
def recommend_recipes(user_input, num_recommendations=5):
    user_input_str = ' '.join(user_input)
    user_input_vector = vectorizer.transform([user_input_str])
    similarity_scores = cosine_similarity(user_input_vector,
ingredient_matrix)
    similarity_scores = similarity_scores.flatten()
    top_indices = similarity_scores.argsort()[-num_recommendations:][::-1]
    return train_df.iloc[top_indices]

# Example usage
user_input = ['chicken', 'tomato', 'rice']
recommendations = recommend_recipes(user_input)

```

```

print("Recommended Recipes:")
print(recommendations[['cuisine', 'ingredients']])

# Evaluate precision and relevance
def evaluate_precision_at_n(user_input, relevant_cuisine, n=5):
    recommendations = recommend_recipes(user_input, num_recommendations=n)
    relevant_recommendations = recommendations[recommendations['cuisine'] ==
relevant_cuisine]
    precision_at_n = len(relevant_recommendations) / n
    return precision_at_n

def evaluate_most_relevant_cuisine(user_input, n=5):
    recommendations = recommend_recipes(user_input, num_recommendations=n)
    cuisine_counts = recommendations['cuisine'].value_counts()
    most_relevant_cuisine = cuisine_counts.idxmax()
    precision_at_n = cuisine_counts.max() / n
    return most_relevant_cuisine, precision_at_n

most_relevant_cuisine, precision = evaluate_most_relevant_cuisine(user_input,
n=5)
print(f'Most Relevant Cuisine: {most_relevant_cuisine}')
print(f'Precision@5 for {most_relevant_cuisine}: {precision:.2f}')

```

Evaluation

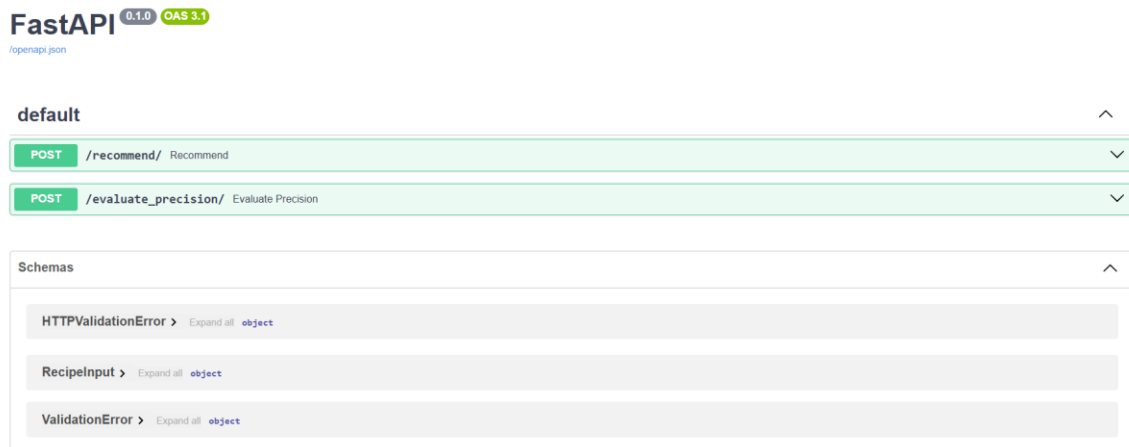
Precision at N: Measures how many of the top-N recommendations match the relevant cuisine.

Most Relevant Cuisine: Identifies the most frequently recommended cuisine among the top-N recommendations.

FastAPI Application

Overview

The FastAPI application serves the trained model to provide real-time recipe recommendations based on user input ingredients.



Application Structure

app.py: Contains the FastAPI application and API endpoints.

Dependencies: pandas, gdown, joblib, fastapi, pydantic, scikit-learn, uvicorn.

API Endpoints

POST /recommend/:

Description: Provides recipe recommendations based on user input ingredients.

Request:

json

Copy code

```
{
  "ingredients": ["chicken", "tomato", "rice"],
  "num_recommendations": 5
}
```

POST

/recommend/ Recommend

Parameters

No parameters

Request body required

```

{
  "ingredients": ["chicken", "tomato", "rice"],
  "num_recommendations": 5
}

```

Response body

```

[
  {
    "cuisine": "mexican",
    "ingredients": ["chicken stock", "hot sauce", "white rice", "chicken", "diced tomatoes", "sweet corn", "garlic"]
  },
  {
    "cuisine": "cajun_creole",
    "ingredients": ["diced onions", "smoked sausage", "shrimp", "crushed tomatoes", "rice", "chicken broth", "chopped celery", "chicken", "cajun seasoning", "carrots"]
  },
  {
    "cuisine": "mexican",
    "ingredients": ["pepper", "diced tomatoes", "onions", "chicken broth", "garlic powder", "salt", "dried basil", "cilantro", "chicken", "tomato sauce", "red pepper", "rotisserie chicken"]
  },
  {
    "cuisine": "mexican",
    "ingredients": ["boneless skinless chicken breasts", "plum tomatoes", "rice", "black beans", "chorizo sausage"]
  },
  {
    "cuisine": "mexican",
    "ingredients": ["chicken and rice soup", "tomato sauce", "tortilla chips", "diced tomatoes", "shredded cheddar cheese"]
  }
]

```

Download

Response: List of recommended recipes with cuisine and ingredients.

POST /evaluate_precision/:

Description: Evaluates the precision of recommendations for a given cuisine.

Request:

json

```

{
  "user_input": ["chicken", "tomato", "rice"],
  "relevant_cuisine": "Italian",
  "n": 5
}

```

Response: Precision at N for the relevant cuisine.

App code –

```
import pandas as pd
import gdown
import joblib
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
from typing import List
from sklearn.metrics.pairwise import cosine_similarity

# Initialize FastAPI app
app = FastAPI()

# Google Drive file IDs (Make sure these files are accessible to 'Anyone with the link')
vectorizer_file_id = '1rcmhf0hxKv1DBjTmCKRnW1FtJiGAaoHP'
dataset_file_id = '1XbiaBkKHmyX5P4eRa05LxYcsSqZfGAF5'

# Download URLs for Google Drive files
vectorizer_download_url =
f'https://drive.google.com/uc?id={vectorizer_file_id}'
dataset_download_url = f'https://drive.google.com/uc?id={dataset_file_id}'

# Paths to save downloaded files
vectorizer_path = 'tfidf_vectorizer.joblib'
dataset_path = 'recipe_dataset.csv'

# Function to download files with error handling
def download_file(url, output_path):
    try:
        gdown.download(url, output_path, quiet=False)
    except Exception as e:
        raise HTTPException(status_code=500, detail=f"Failed to download {output_path}: {str(e)}")

# Download the necessary files
download_file(vectorizer_download_url, vectorizer_path)
download_file(dataset_download_url, dataset_path)

# Load the pre-trained TF-IDF vectorizer model
try:
    vectorizer = joblib.load(vectorizer_path)
except Exception as e:
    raise HTTPException(status_code=500, detail=f"Failed to load vectorizer: {str(e)}")

# Load the dataset
try:
    train_df = pd.read_csv(dataset_path)
```



```

except Exception as e:
    raise HTTPException(status_code=500, detail=f"Failed to load dataset: {str(e)}")

# Ensure the 'cleaned_ingredients' column exists
if 'cleaned_ingredients' not in train_df.columns:
    raise HTTPException(status_code=500, detail="Column 'cleaned_ingredients' not found in dataset.")

# Vectorize the ingredients using the loaded vectorizer
ingredient_matrix = vectorizer.transform(train_df['cleaned_ingredients'])

# Define input model for FastAPI
class RecipeInput(BaseModel):
    ingredients: List[str]
    num_recommendations: int = 5

# Recommendation function
def recommend_recipes(user_input, num_recommendations=5):
    user_input_str = ' '.join(user_input)
    user_input_vector = vectorizer.transform([user_input_str])
    similarity_scores = cosine_similarity(user_input_vector,
ingredient_matrix)
    similarity_scores = similarity_scores.flatten()
    top_indices = similarity_scores.argsort()[-num_recommendations:][::-1]
    return train_df.iloc[top_indices]

# API Endpoint for recommending recipes
@app.post("/recommend/")
def recommend(data: RecipeInput):
    try:
        recommendations = recommend_recipes(data.ingredients,
data.num_recommendations)
        return recommendations[['cuisine',
'ingredients']].to_dict(orient='records')
    except Exception as e:
        raise HTTPException(status_code=500, detail=f"Error during
recommendation: {str(e)}")

# Precision evaluation function
@app.post("/evaluate_precision/")
def evaluate_precision(user_input: List[str], relevant_cuisine: str, n: int = 5):
    try:
        recommendations = recommend_recipes(user_input, num_recommendations=n)
        relevant_recommendations = recommendations[recommendations['cuisine']
== relevant_cuisine]
        precision_at_n = len(relevant_recommendations) / n

```


Deployment Setup

Railway Deployment - [FastAPI - Swagger UI \(railway.app\)](#)

Preparation:

Ensure the following files are present in the project directory:

- Procfile
- requirements.txt
- app.py

Procfile:

```
web: uvicorn app:app --host 0.0.0.0 --port $PORT
```

requirements.txt:

```
fastapi
uvicorn
scikit-learn
pandas
gdown
joblib
```

Deployment Steps

- upload locally tested files to git hub.
- connect railway acc to GitHub
- choose repository that we need to deploy
- deploy that application from railway
- open the application

CI/CD Pipeline Documentation

Overview

The CI/CD pipeline automates the testing, building, and deployment process for the FastAPI application. This setup ensures that code changes are automatically tested and deployed to Railway without manual intervention.

CI/CD Pipeline Setup

GitHub Actions Workflow

Purpose: This GitHub Actions workflow file defines the steps for testing, building, and deploying the application to Railway.

Contents-

yaml

name: Deploy FastAPI App

on:

push:

branches:

- master

jobs:

build:

runs-on: ubuntu-latest

steps:

- name: Checkout code

uses: actions/checkout@v2

with:

Checkout the code from the repository

- name: Set up Python

uses: actions/setup-python@v2

with:

python-version: '3.9'

```

# Set up Python environment version 3.9

- name: Install dependencies
  run: |
    pip install -r requirements.txt
    # Install the dependencies specified in requirements.txt

- name: Run tests
  run: |
    # Add commands to run tests if applicable
    echo "No tests found."
    # Placeholder for test execution; customize as needed

- name: Deploy to Railway
  env:
    RAILWAY_API_KEY: ${ secrets.RAILWAY_API_KEY }
  run: |
    railway login --token $RAILWAY_API_KEY
    railway up
    # Log in to Railway CLI and deploy the application

```

Explanation of Steps

Checkout Code: Retrieves the latest code from the GitHub repository.

Set up Python: Sets up the Python environment with the specified version (3.9).

Install Dependencies: Installs the Python packages listed in requirements.txt.

Run Tests: Executes tests to ensure the application works as expected. If no tests are present, a placeholder message is displayed.

Deploy to Railway: Logs in to Railway CLI using the API key stored in GitHub secrets and deploys the application.

2. Environment Variables

RAILWAY_API_KEY: This secret is used to authenticate and deploy the application to Railway. It is stored securely in GitHub secrets.

To Add the Secret:

Navigate to your GitHub repository.

Go to Settings > Secrets and variables > Actions.

Click New repository secret.

Add RAILWAY_API_KEY as the name and your Railway API key as the value.

Deployment Process

Push Changes: When changes are pushed to the master branch, the workflow is triggered.

Checkout Code: The code is checked out from the repository.

Set up Python Environment: The appropriate Python version is set up.

Install Dependencies: Required Python packages are installed.

Run Tests: Any tests are executed (if defined).

Deploy Application: The application is deployed to Railway using the Railway CLI.

Online Testing

After deployment, Railway provides a URL to access the application. You can find the URL in the Railway dashboard under your project details. The URL typically follows the format:

<https://web-production-a2847.up.railway.app/docs#>

Conclusion

Summary

This documentation covers the model training process, FastAPI application setup, deployment to railway, and CI/CD pipeline configuration for the recipe recommendation system.